



a Head-by-head robustness: where does each model succeed or struggle?



b Whole-cascade robustness: were all applicable Heads correct?




c 1 Select on Train CV only


H3 encoder

	C-300M	C-600M	C-6B
ESM-2 650M	0.9949 ±0.0016	0.9950 ±0.0015	0.9968 ±0.0010
ESM-2 3B	0.9958 ±0.0016	0.9959 ±0.0016	0.9976 ±0.0010 ★
ESM-C 6B	0.9953 ±0.0017	0.9954 ±0.0015	0.9971 ±0.0009

2 Freeze Head assignment



3 Check neighbours — do not rerank



0/4 inferiority warnings vs all-6B

$S = 0.60 \text{ H1 AP} + 0.30 \text{ H2 AP} + 0.10 \text{ H3 known macro-F1}$
cell = five-fold mean $S \pm \text{fold SE}$

No warning ≠ equivalence · Test accessed = 0
External confirmation required · V0 unchanged